



Czech & Slovak



KCD

GROWING CLOUD NATIVE TOGETHER

May 21 - 22, 2026 | PRAGUE

I'm speaking at KCD Czech & Slovak 2026

Craft your home AI lab with Ollama + Open WebUI



Petr Hanák

About me... boring

Teacher (Mathematics+Informatics) >>

>> Microsoft Trainer (Active Directory, Windows Servers...) >>

>> Windows Admin >>

>> DevOps (Cloud Native, K8S, ci/cd) >>

>> DevOps@SAP (SAP Labs Brno)

petr.hanak@email.cz

www.linkedin.com/in/petrhanak/

Why Local LLM (answered by Opus 4.7)

- >> PRIVACY & SECURITY (Your data never leaves your infrastructure)
- >> COST (one time HW price x tokens, predictable)
- >> OFFLINE (without internet + latency)
- >> FULL CONTROL + CUSTOMIZATION (no restriction + no rate limits)
- >> LEARNING + EXPERIMENTATION

Why not local (answered by Sonnet 4.6)

- >> Need best quality
- >> Low/occasional usage
- >> Don't want infrastructure
- >> Size (forgets context after ~8k tokens)
- >> Not creative, no real-time knowledge
- >> You need serious HW (GPU like RTX 3060+)
- >> Setup overhead (why we are here)

Hybrid?

	Local LLM	Cloud API
Quality (8B model)	☆☆☆	☆☆☆☆☆
Quality (32B model)	☆☆☆☆	☆☆☆☆☆
Setup difficulty	Hard ✗	Easy ✓
Maintenance	You ✗	Provider ✓
Scaling	Limited ✗	Infinite ✓
Privacy	Perfect ✓	Good ⚠
Cost (low volume)	Expensive ✗	Cheap ✓
Cost (high volume)	Cheap ✓	Expensive ✗
Offline use	Yes ✓	No ✗
Latest knowledge	No ✗	Yes ✓
Reliability	Your problem	SLA ✓
Tool/function calling	Unreliable ⚠	Reliable ✓
Context window	8-32k ⚠	1M tokens ✓

How to run local LLM? SW+HW

HW?

>> Experiments? just CPU. Deep dive? capable GPU

>> Virtual (run.pod.io starts at ~\$0.14/hr) or physical (RTX3060 12GB - € 300)

SW?

>> User Interface - CLI or Open WebUI

>> Runtime - **Ollama** (or LM Studio, llama.cpp)

Usage?

>> Chat, coding, agents, images ...

Goal of the workshop?

Having running local LLM model accessible via command line, API and GUI and ready for further experiments like “vibe coding”.

Our setup? (your laptop)+runpod.io

For Lab A - your laptop (or your VM) or cpu-xx pod in runpod.io

>> Web terminal

For Lab B - pod gpu-xx with Ollama+Open WebUI accessible via 3 URL's

>> Web terminal

>> Ollama API

>> Open-WebUI

Let's start...

<https://github.com/petrhanakcz/ollama-lab>

Open

../StudentFiles/student-xx.md file

../materials/Craft_home_LAB_with_Ollama.md